

Assessing the Quality of Activity

by clearwriter, February 2016

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit, it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information on the original research is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>) (see the section on the Weight Lifting Exercise Dataset).

Data Source

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)

The test data are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>). If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

Load Libraries and Prepare Datasets

Load libraries first.

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(parallel)  
library(doParallel)
```

```
## Loading required package: foreach
```

```
## Loading required package: iterators
```

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
##     margin
```

```
library(rpart)  
library(rpart.plot)  
library(RColorBrewer)  
library(sjPlot)  
library(knitr)  
library(captioner)  
library(doMC)  
require(data.table)
```

```
## Loading required package: data.table
```

```
set.seed(1234)
```

Download testing and training data to your working directory.

```
## Load training data.  
url <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"  
training <- fread(url)  
url <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"  
testing <- fread(url)  
dim(training); dim(testing);
```

```
## [1] 19622   160
```

```
## [1] 20 160
```

Identify Predictor Candidates

Which variables in the test dataset have zero NAs? Use this tip: finding columns with all missing values in `r`.

Belt, arm, dumbbell, and forearm variables that do not have any missing values in the test dataset will be predictor candidates.

```
isAnyMissing <- sapply(testing, function (x) any(is.na(x) | x == ""))
isPredictor <- !isAnyMissing & grepl("belt|^(fore)arm|dumbbell|forearm", names(isAnyMissing))
predCandidates <- names(isAnyMissing)[isPredictor]
predCandidates
```

```
## [1] "roll_belt" "pitch_belt" "yaw_belt"
## [4] "total_accel_belt" "gyros_belt_x" "gyros_belt_y"
## [7] "gyros_belt_z" "accel_belt_x" "accel_belt_y"
## [10] "accel_belt_z" "magnet_belt_x" "magnet_belt_y"
## [13] "magnet_belt_z" "roll_arm" "pitch_arm"
## [16] "yaw_arm" "total_accel_arm" "gyros_arm_x"
## [19] "gyros_arm_y" "gyros_arm_z" "accel_arm_x"
## [22] "accel_arm_y" "accel_arm_z" "magnet_arm_x"
## [25] "magnet_arm_y" "magnet_arm_z" "roll_dumbbell"
## [28] "pitch_dumbbell" "yaw_dumbbell" "total_accel_dumbbell"
## [31] "gyros_dumbbell_x" "gyros_dumbbell_y" "gyros_dumbbell_z"
## [34] "accel_dumbbell_x" "accel_dumbbell_y" "accel_dumbbell_z"
## [37] "magnet_dumbbell_x" "magnet_dumbbell_y" "magnet_dumbbell_z"
## [40] "roll_forearm" "pitch_forearm" "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x" "gyros_forearm_y"
## [46] "gyros_forearm_z" "accel_forearm_x" "accel_forearm_y"
## [49] "accel_forearm_z" "magnet_forearm_x" "magnet_forearm_y"
## [52] "magnet_forearm_z"
```

Next, we want to subset the primary dataset to include only the predictor candidates and the outcome variable, `classe`.

```
varToInclude <- c("classe", predCandidates)
training <- training[, varToInclude, with=FALSE]
dim(training)
```

```
## [1] 19622 53
```

```
names(training)
```

```
## [1] "classe" "roll_belt" "pitch_belt"
## [4] "yaw_belt" "total_accel_belt" "gyros_belt_x"
## [7] "gyros_belt_y" "gyros_belt_z" "accel_belt_x"
## [10] "accel_belt_y" "accel_belt_z" "magnet_belt_x"
## [13] "magnet_belt_y" "magnet_belt_z" "roll_arm"
## [16] "pitch_arm" "yaw_arm" "total_accel_arm"
## [19] "gyros_arm_x" "gyros_arm_y" "gyros_arm_z"
## [22] "accel_arm_x" "accel_arm_y" "accel_arm_z"
## [25] "magnet_arm_x" "magnet_arm_y" "magnet_arm_z"
## [28] "roll_dumbbell" "pitch_dumbbell" "yaw_dumbbell"
## [31] "total_accel_dumbbell" "gyros_dumbbell_x" "gyros_dumbbell_y"
## [34] "gyros_dumbbell_z" "accel_dumbbell_x" "accel_dumbbell_y"
## [37] "accel_dumbbell_z" "magnet_dumbbell_x" "magnet_dumbbell_y"
## [40] "magnet_dumbbell_z" "roll_forearm" "pitch_forearm"
## [43] "yaw_forearm" "total_accel_forearm" "gyros_forearm_x"
## [46] "gyros_forearm_y" "gyros_forearm_z" "accel_forearm_x"
## [49] "accel_forearm_y" "accel_forearm_z" "magnet_forearm_x"
## [52] "magnet_forearm_y" "magnet_forearm_z"
```

And then we convert classe into a factor.

```
training <- training[, classe := factor(training[, classe])]
training[, .N, classe]
```

```
##   classe    N
## 1:     A 5580
## 2:     B 3797
## 3:     C 3422
## 4:     D 3216
## 5:     E 3607
```

As we've learned, we split the dataset into 60/40 training/probing.

```
inTrain <- createDataPartition(training$classe, p=0.6)
DTrain <- training[inTrain[[1]]]
DProbe <- training[-inTrain[[1]]]
```

We preprocess the prediction variables by centering and scaling.

```
X <- DTrain[, predCandidates, with=FALSE]
preProc <- preProcess(X)
preProc
```

```
## Created from 11776 samples and 52 variables
##
## Pre-processing:
##   - centered (52)
##   - ignored (0)
##   - scaled (52)
```

```
XCS <- predict(preProc, X)
DTrainCS <- data.table(data.frame(classe = DTrain[, classe], XCS))
```

And then apply the centering and scaling to our probing dataset.

```
X <- DProbe[, predCandidates, with=FALSE]
XCS <- predict(preProc, X)
DProbeCS <- data.table(data.frame(classe = DProbe[, classe], XCS))
```

We also need to check for near zero variance.

```
nzv <- nearZeroVar(DTrainCS, saveMetrics=TRUE)
if (any(nzv$nzv)) nzv else message("No variables with near zero variance")
```

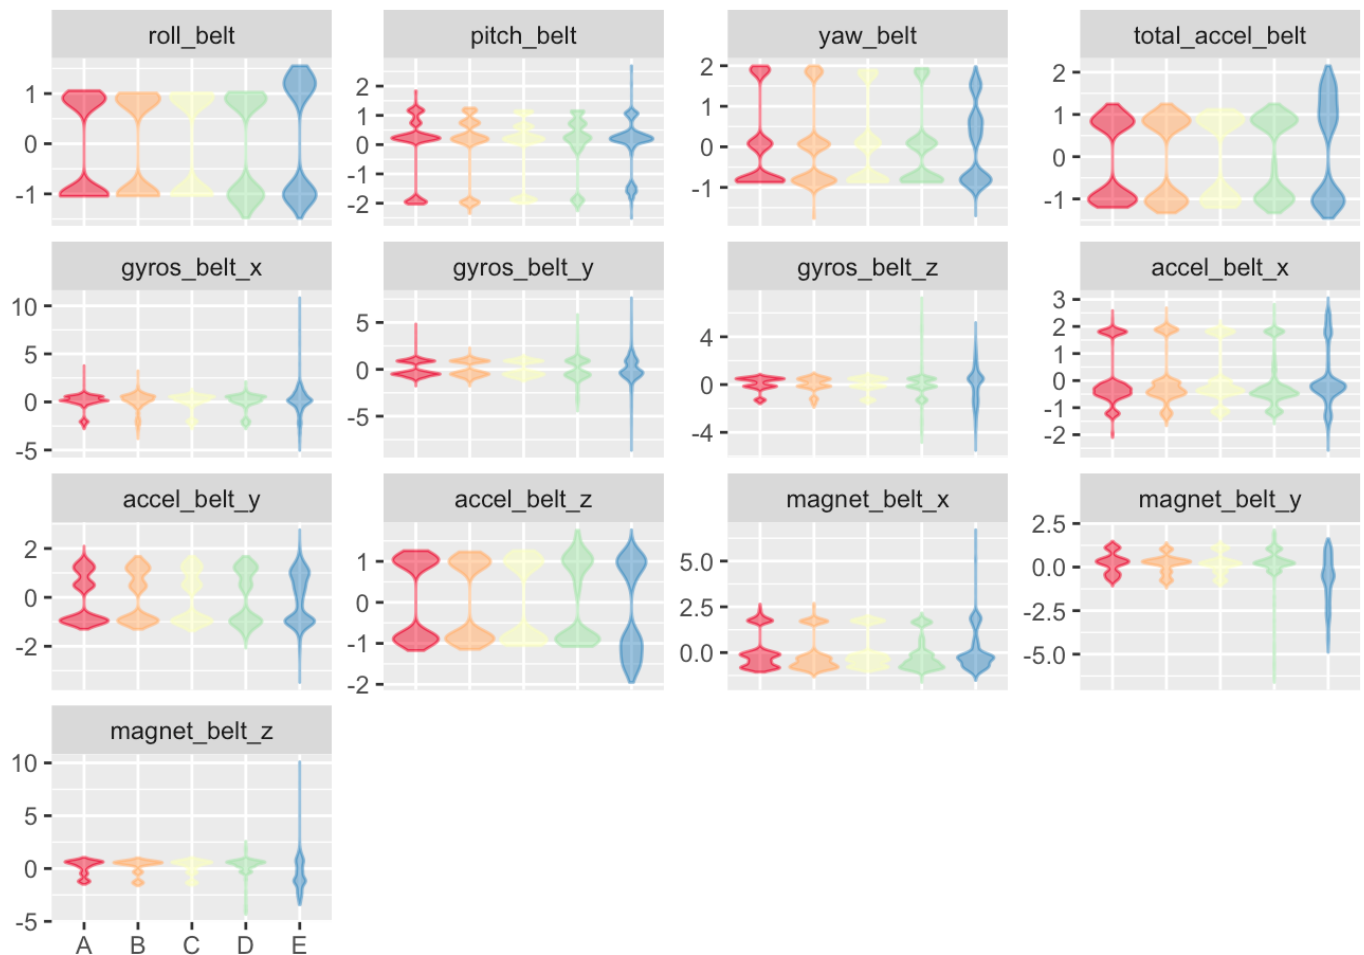
```
## No variables with near zero variance
```

Now, let's examine our groups of prediction variables.

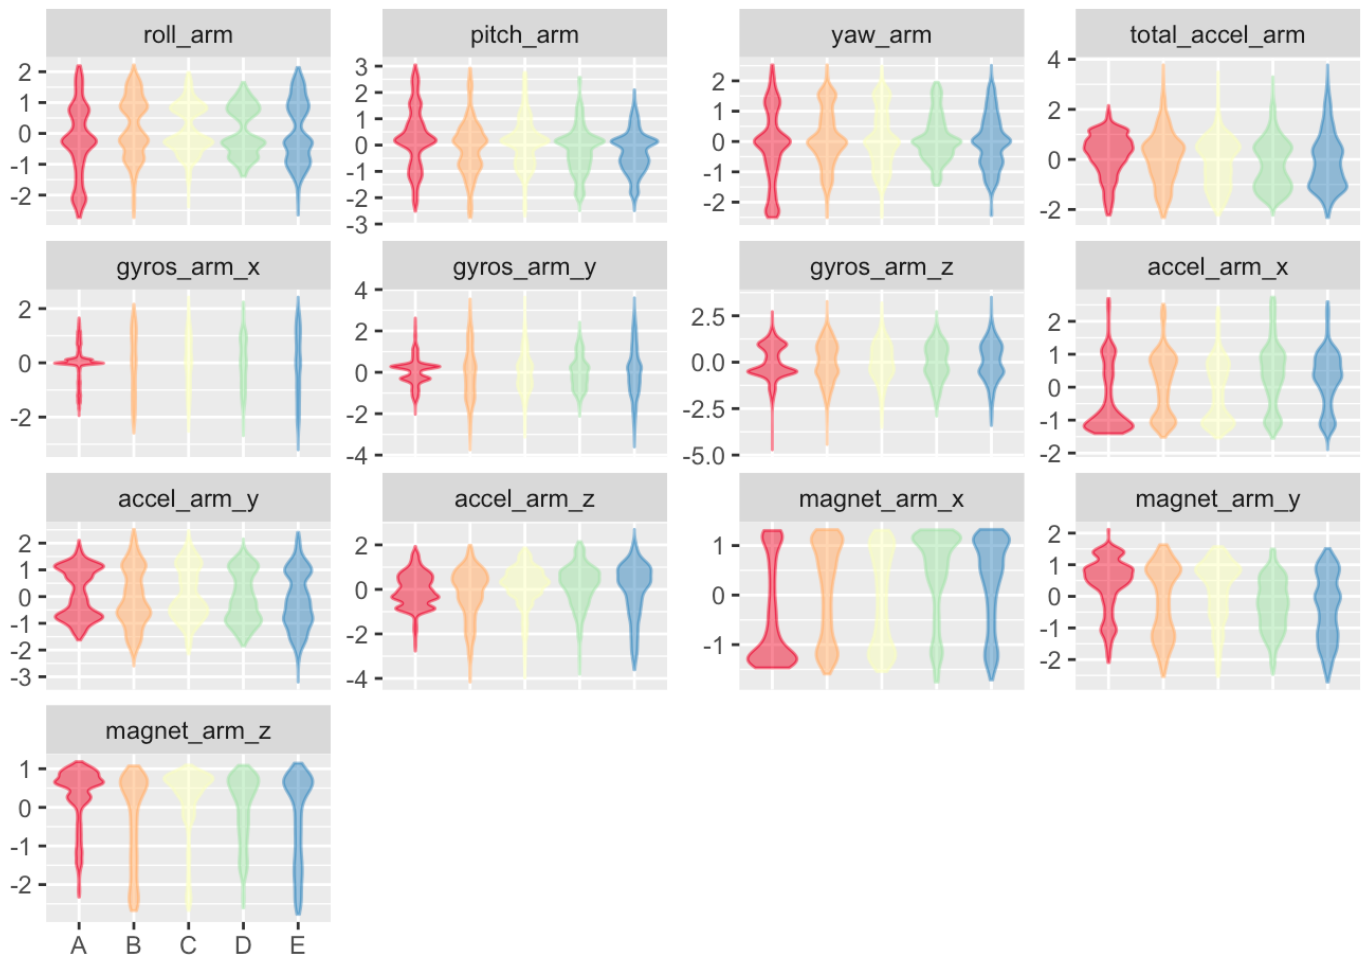
```
histGroup <- function (data, regex) {
  col <- grep(regex, names(data))
  col <- c(col, which(names(data) == "classe"))
  library(reshape2)
  n <- nrow(data)
  DMelted <- melt(data[, col, with=FALSE][, rownum := seq(1, n)], id.vars=c("rownum", "classe"))
  library(ggplot2)
  ggplot(DMelted, aes(x=classe, y=value)) +
    geom_violin(aes(color=classe, fill=classe), alpha=1/2) +
  #   geom_jitter(aes(color=classe, fill=classe), alpha=1/10) +
  #   geom_smooth(aes(group=1), method="gam", color="black", alpha=1/2, size=2) +
  facet_wrap(~ variable, scale="free_y") +
  scale_color_brewer(palette="Spectral") +
  scale_fill_brewer(palette="Spectral") +
  labs(x="", y="") +
  theme(legend.position="none")
}
histGroup(DTrainCS, "belt")
```

```
##  
## Attaching package: 'reshape2'
```

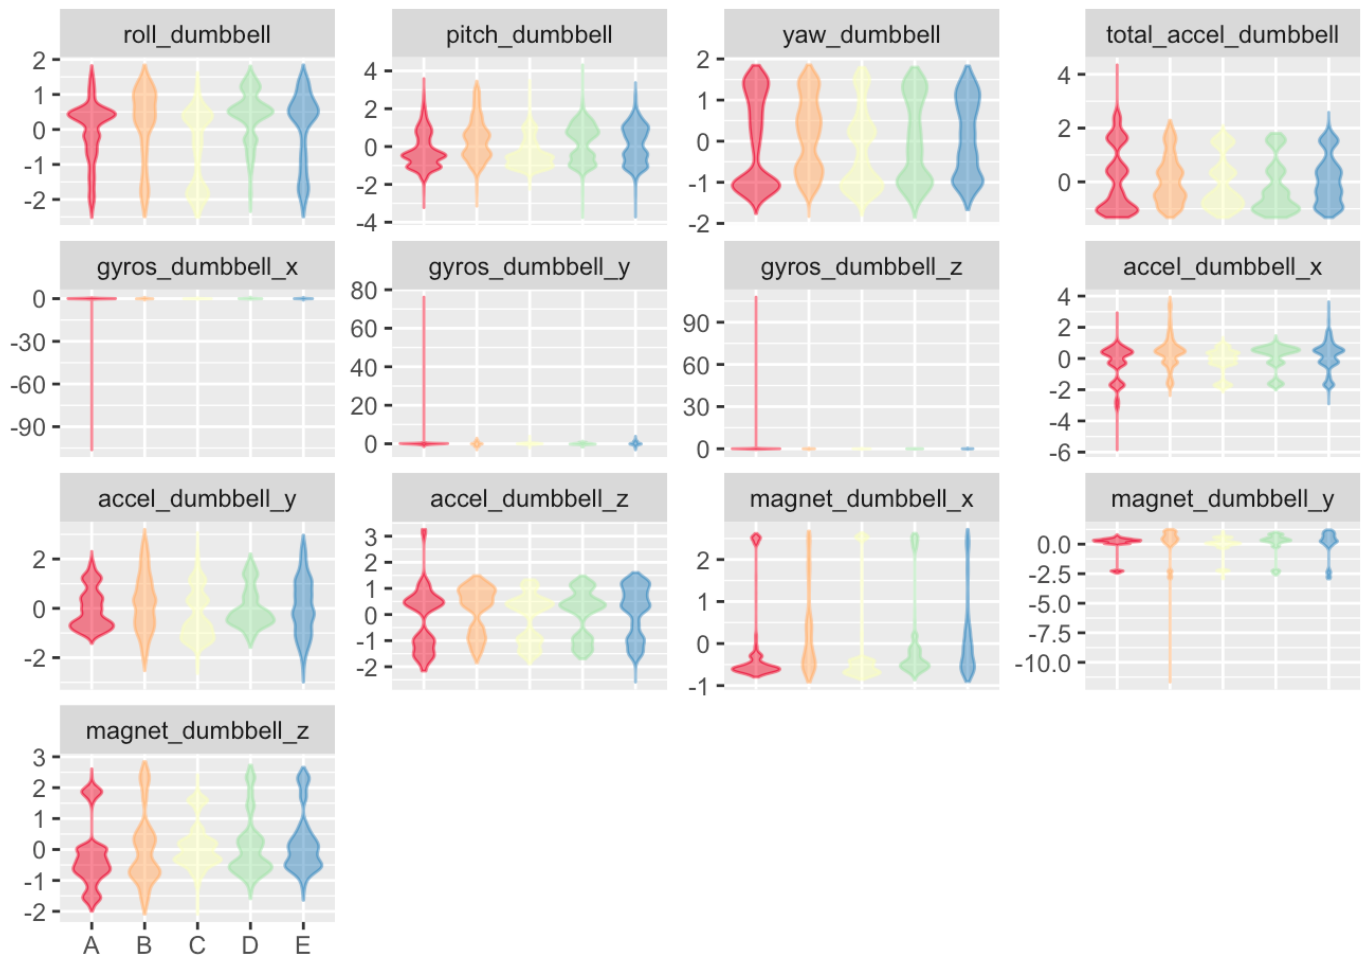
```
## The following objects are masked from 'package:data.table':  
##  
## dcast, melt
```



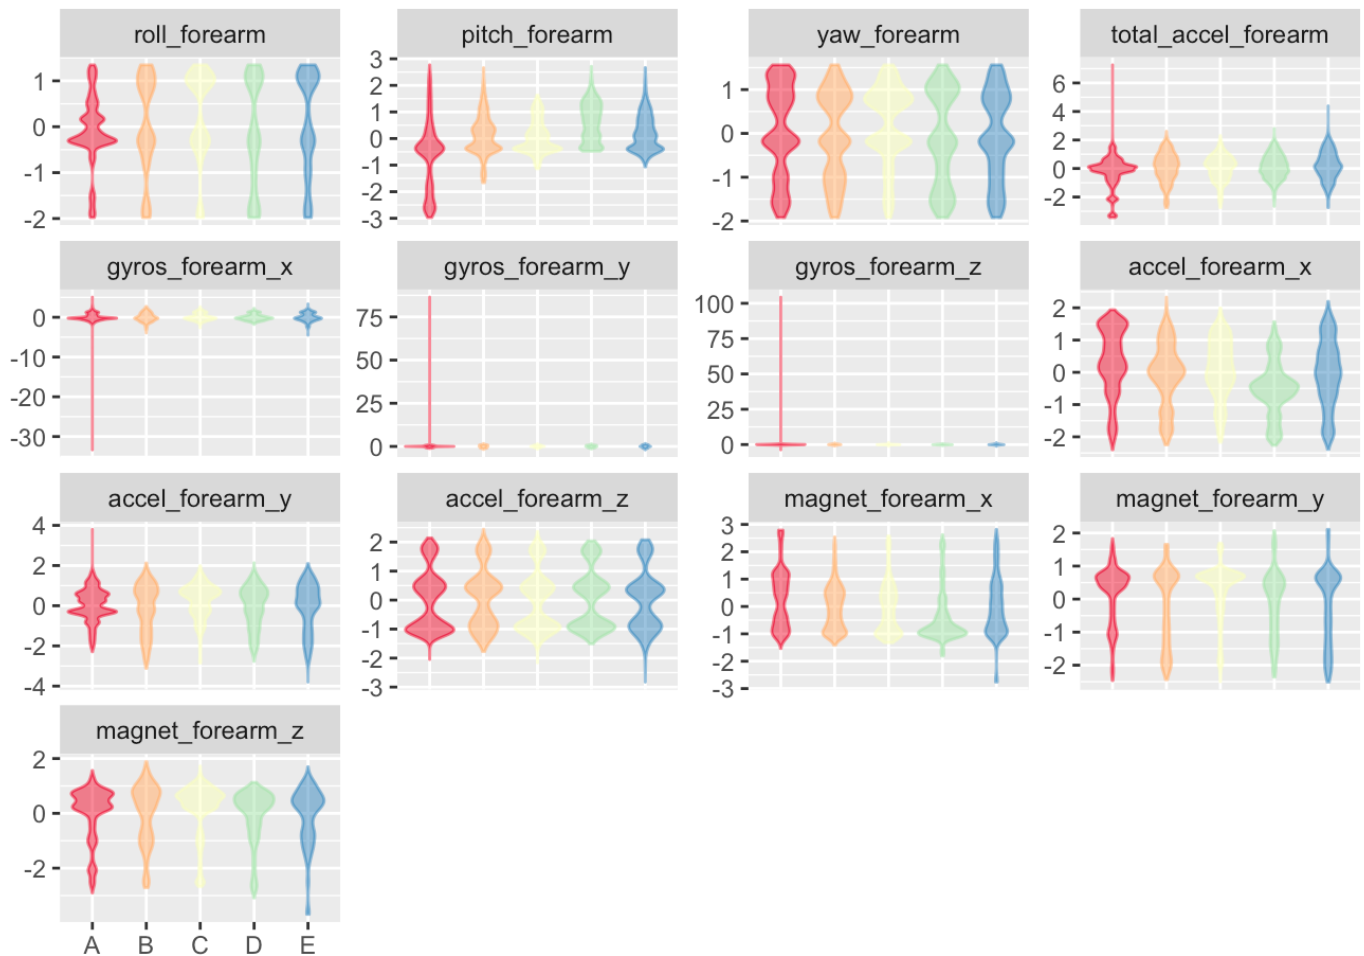
```
histGroup(DTrainCS, "[^(fore)]arm")
```



```
histGroup(DTrainCS, "dumbbell")
```



```
histGroup(DTrainCS, "forearm")
```

Training a Prediction Model

Using a random forest, the out-of-sample error should be small. We'll estimate the error using 40% probing sample.

Set up the parallel clusters.

```
cl <- makeCluster(detectCores() - 1)
registerDoParallel(cl)
```

Set the control parameters.

```
ctrl <- trainControl(classProbs=TRUE,
                      savePredictions=TRUE,
                      allowParallel=TRUE)
```

And fit our model over the training parameters. Note: this takes a while.

```
method <- "rf"
system.time(trainingModel <- train(classe ~ ., data=DTrainCS, method=method))
```

```
##      user    system elapsed
## 34.518      1.085 1934.071
```

```
stopCluster(cl)
```

Evaluate the Training Model

```
trainingModel
```

```
## Random Forest
##
## 11776 samples
##      52 predictor
##      5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
##      mtry  Accuracy   Kappa      Accuracy SD   Kappa SD
##      2     0.9861383  0.9824602  0.002500503    0.003172630
##     27     0.9857348  0.9819517  0.002154335    0.002727942
##     52     0.9771254  0.9710582  0.005047104    0.006397850
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 2.
```

```
hat <- predict(trainingModel, DTrainCS)
confusionMatrix(hat, DTrain[, classe])
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 3348    0    0    0    0
##           B    0 2279    0    0    0
##           C    0    0 2054    0    0
##           D    0    0    0 1930    0
##           E    0    0    0    0 2165
##
## Overall Statistics
##
##           Accuracy : 1
##           95% CI : (0.9997, 1)
##           No Information Rate : 0.2843
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 1
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000   1.0000   1.0000   1.0000   1.0000
## Specificity      1.0000   1.0000   1.0000   1.0000   1.0000
## Pos Pred Value   1.0000   1.0000   1.0000   1.0000   1.0000
## Neg Pred Value   1.0000   1.0000   1.0000   1.0000   1.0000
## Prevalence       0.2843   0.1935   0.1744   0.1639   0.1838
## Detection Rate   0.2843   0.1935   0.1744   0.1639   0.1838
## Detection Prevalence 0.2843   0.1935   0.1744   0.1639   0.1838
## Balanced Accuracy 1.0000   1.0000   1.0000   1.0000   1.0000
```

Evaluate the Model Using the Probing Dataset

```
hat <- predict(trainingModel, DProbeCS)
confusionMatrix(hat, DProbeCS[, classe])
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2231   11    0    0    0
##           B    1 1504   11    0    0
##           C    0    3 1353   31    3
##           D    0    0    4 1254    2
##           E    0    0    0    1 1437
##
## Overall Statistics
##
##           Accuracy : 0.9915
##           95% CI : (0.9892, 0.9934)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9892
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9996   0.9908   0.9890   0.9751   0.9965
## Specificity      0.9980   0.9981   0.9943   0.9991   0.9998
## Pos Pred Value   0.9951   0.9921   0.9734   0.9952   0.9993
## Neg Pred Value   0.9998   0.9978   0.9977   0.9951   0.9992
## Prevalence       0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate   0.2843   0.1917   0.1724   0.1598   0.1832
## Detection Prevalence 0.2858   0.1932   0.1772   0.1606   0.1833
## Balanced Accuracy 0.9988   0.9944   0.9917   0.9871   0.9982
```

Final Model

```
varImp(trainingModel)
```

```
## rf variable importance
##
##   only 20 most important variables shown (out of 52)
##
##               Overall
## roll_belt      100.00
## yaw_belt       78.03
## magnet_dumbbell_z 63.91
## magnet_dumbbell_y 62.02
## pitch_forearm  61.83
## pitch_belt     59.26
## magnet_dumbbell_x 51.14
## roll_forearm   50.20
## magnet_belt_z  43.52
## roll_dumbbell  42.57
## accel_dumbbell_y 42.52
## magnet_belt_y  42.42
## accel_belt_z   41.28
## accel_dumbbell_z 37.17
## roll_arm       34.57
## accel_forearm_x 30.70
## accel_dumbbell_x 29.08
## yaw_dumbbell   28.79
## gyros_dumbbell_y 27.33
## magnet_forearm_z 27.25
```

```
trainingModel$finalModel
```

```
##
## Call:
##   randomForest(x = x, y = y, mtry = param$mtry)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 2
##
##               OOB estimate of  error rate: 0.85%
## Confusion matrix:
##      A    B    C    D    E  class.error
## A 3346    2    0    0    0 0.0005973716
## B   19 2255    5    0    0 0.0105309346
## C    1   23 2027    3    0 0.0131450828
## D    0    0   35 1893    2 0.0191709845
## E    0    0    5    5 2155 0.0046189376
```

We have an estimated error rate of less than 1%. Excellent. We'll save this training model for later.

```
save(trainingModel, file="trainingModel.RData")
```

Predictions

Load the training model.

```
load(file="trainingModel.RData", verbose=TRUE)
```

```
## Loading objects:  
##   trainingModel
```

Predict and evaluate.

```
DTestCS <- predict(preProc, testing[, predCandidates, with=FALSE])  
hat <- predict(trainingModel, DTestCS)  
testing <- cbind(hat , testing)  
subset(testing, select=names(testing)[grep("belt|^[^fore]arm|dumbbell|forearm", names(testing), invert=TRUE)])
```

##	hat	V1	user_name	raw_timestamp_part_1	raw_timestamp_part_2
## 1:	B	1	pedro	1323095002	868349
## 2:	A	2	jeremy	1322673067	778725
## 3:	B	3	jeremy	1322673075	342967
## 4:	A	4	adelmo	1322832789	560311
## 5:	A	5	eurico	1322489635	814776
## 6:	E	6	jeremy	1322673149	510661
## 7:	D	7	jeremy	1322673128	766645
## 8:	B	8	jeremy	1322673076	54671
## 9:	A	9	carlitos	1323084240	916313
## 10:	A	10	charles	1322837822	384285
## 11:	B	11	carlitos	1323084277	36553
## 12:	C	12	jeremy	1322673101	442731
## 13:	B	13	eurico	1322489661	298656
## 14:	A	14	jeremy	1322673043	178652
## 15:	E	15	jeremy	1322673156	550750
## 16:	E	16	eurico	1322489713	706637
## 17:	A	17	pedro	1323094971	920315
## 18:	B	18	carlitos	1323084285	176314
## 19:	B	19	pedro	1323094999	828379
## 20:	B	20	eurico	1322489658	106658

##	cvtd_timestamp	new_window	num_window	problem_id
## 1:	05/12/2011 14:23	no	74	1
## 2:	30/11/2011 17:11	no	431	2
## 3:	30/11/2011 17:11	no	439	3
## 4:	02/12/2011 13:33	no	194	4
## 5:	28/11/2011 14:13	no	235	5
## 6:	30/11/2011 17:12	no	504	6
## 7:	30/11/2011 17:12	no	485	7
## 8:	30/11/2011 17:11	no	440	8
## 9:	05/12/2011 11:24	no	323	9
## 10:	02/12/2011 14:57	no	664	10
## 11:	05/12/2011 11:24	no	859	11
## 12:	30/11/2011 17:11	no	461	12
## 13:	28/11/2011 14:14	no	257	13
## 14:	30/11/2011 17:10	no	408	14
## 15:	30/11/2011 17:12	no	779	15
## 16:	28/11/2011 14:15	no	302	16
## 17:	05/12/2011 14:22	no	48	17
## 18:	05/12/2011 11:24	no	361	18
## 19:	05/12/2011 14:23	no	72	19
## 20:	28/11/2011 14:14	no	255	20